

August 2021 Inflation Forecasting

Kyle Mouly

September 13, 2021

Contents

Setup	1
Loading Data	1
Stationarity and Seasonality	2
Univariate Models	4
Model Selection	4
Forecasting	7
Official Results	8

Setup

For this project I will be forecasting non-seasonally adjusted year over year CPI for all Urban Consumers for the month of August 2021 (FRED series id: CPIAUCNS). The data series will be segmented so as to start after September 1987 to account for the change in the Federal Reserve's attitude towards inflation following Paul Volcker's tenure as chair; a change which is clear when long term inflation trends are examined. The loss of data should not be concerning given that there remains over 30 years of monthly data for use in model training.

Loading Data

```
library(tidyverse)
library(fable)
library(fabletools)
library(tsibble)
library(fredr)
library(feasts)
library(lubridate)


CPI <- fredr(series_id = "CPIAUCNS",
             observation_start = as.Date("1987-09-30"),
             observation_end = as.Date("2021-07-30"),
             units = "pc1") %>%
  mutate(date = yearmonth(date)) %>%
  select(-c("realtime_start", "realtime_end", "series_id")) %>%
  rename(pc1_cpi = value) %>%
  as_tsibble(index = date)

CPI %>% autoplot(pc1_cpi)+
  labs(title = "U.S. CPI All Urban Consumers",
```

```
subtitle = "All Items U.S. City Average",  
y = "YOY Percent Change CPI")
```

U.S. CPI All Urban Consumers


All Items U.S. City Average



Plot does not indicate any obvious stationary, trending, seasonal behavior; formal unit root tests will be necessary to determine what transformations if any are necessary.

Stationarity and Seasonality

```
CPI %>%  
  gg_tsdisplay(pc1_cpi, "partial")
```



```
CPI %>%
  features(pc1_cpi,
    features = list(
      unitroot_kpss,
      unitroot_ndiffs,
      unitroot_nsdiffs))

## # A tibble: 1 x 4
##   kpss_stat kpss_pvalue ndiffs nsdiffs
##   <dbl>     <dbl> <int>  <int>
## 1      2.32      0.01     1      0

CPI %>%
  mutate(dpc1_cpi = difference(pc1_cpi)) %>%
  features(dpc1_cpi, unitroot_kpss)


## # A tibble: 1 x 2
##   kpss_stat kpss_pvalue
##   <dbl>     <dbl>
## 1    0.0775      0.1
```

As per the KPSS test the data is non-stationary and requires a first difference to become stationary.

Univariate Models

Model Selection

```
CPI %>%
  slice(1:(n()-12)) %>%
  mutate(dpc1_cpi = difference(pc1_cpi)) %>%
  na.omit() %>%
  gg_tsdisplay(dpc1_cpi, plot_type = "partial")
```



ACF suggests an $ARIMA(0,1,1)(0,0,2)[12]$ candidate model. The PACF suggests an $ARIMA(2,1,0)(0,0,2)[12]$ candidate model.

```
HCV.fit <- CPI %>%
  slice(1:(n()-12)) %>%
  model(ARIMA011 = ARIMA(pc1_cpi ~ pdq(0,1,1) + PDQ(0,0,2)),
        ARIMA210 = ARIMA(pc1_cpi ~ pdq(2,1,0) + PDQ(0,0,2)),
        arimaAICc = ARIMA(pc1_cpi, ic = "aicc", stepwise = FALSE, approximation = FALSE),
        arimaBIC = ARIMA(pc1_cpi, ic = "bic", stepwise = FALSE, approximation = FALSE),
        etsAICc = ETS(pc1_cpi, ic = "aicc"),
        etsBIC = ETS(pc1_cpi, ic = "bic"),
        naive = NAIVE(pc1_cpi)
  )
```

```
HCV.fit %>% pivot_longer(everything(), names_to = "Model", values_to = "Order")
```

```
## # A mable: 7 x 2
```

```
## # Key:      Model [7]
##   Model                                Order
##   <chr>                                <model>
## 1 ARIMA011      <ARIMA(0,1,1)(0,0,2)[12]>
## 2 ARIMA210      <ARIMA(2,1,0)(0,0,2)[12]>
## 3 arimaAICc     <ARIMA(2,1,0)(2,0,1)[12] w/ drift>
## 4 arimaBIC      <ARIMA(0,1,1)(0,0,2)[12]>
## 5 etsAICc       <ETS(A,Ad,N)>
## 6 etsBIC        <ETS(A,N,N)>
## 7 naive         <NAIVE>


HCV.fit %>% glance() %>% arrange(AICc) %>% select(.model:BIC)
```

```
## # A tibble: 7 x 6
##   .model      sigma2 log_lik      AIC   AICc     BIC
##   <chr>        <dbl>   <dbl>  <dbl>  <dbl>  <dbl>
## 1 arimaAICc  0.0597   -12.0   38.1   38.4   65.9
## 2 ARIMA210   0.0611   -16.4   42.8   43.0   62.7
## 3 ARIMA011   0.0617   -18.2   44.3   44.4   60.2
## 4 arimaBIC   0.0617   -18.2   44.3   44.4   60.2
## 5 etsAICc    0.141    -791.  1595.  1595.  1619.
## 6 etsBIC     0.145    -799.  1603.  1603.  1615.
## 7 naive      0.145      NA     NA     NA     NA
```

```
bind_rows(
  HCV.fit %>% accuracy(),
  HCV.fit %>% forecast(h = 12) %>% accuracy(CPI)
) %>%
  arrange(RMSE)
```

```
## # A tibble: 14 x 10
##   .model      .type      ME  RMSE  MAE  MPE  MAPE  MASE  RMSSE      ACF1
##   <chr>        <chr>    <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>    <dbl>
## 1 arimaAICc  Training -0.00410 0.242 0.176 -3.13 22.3 0.158 0.157 0.000269
## 2 ARIMA210   Training -0.0314 0.246 0.179 -3.23 21.8 0.161 0.159 -0.0110
## 3 ARIMA011   Training -0.0312 0.247 0.178 -3.21 21.1 0.160 0.160 0.0313
## 4 arimaBIC   Training -0.0312 0.247 0.178 -3.21 21.1 0.160 0.160 0.0313
## 5 etsAICc    Training -0.00142 0.373 0.264 1.31 30.1 0.237 0.242 0.249
## 6 etsBIC     Training -0.00853 0.380 0.262 -2.77 25.0 0.235 0.247 0.387
## 7 naive      Training -0.00855 0.381 0.263 -2.78 25.1 0.236 0.247 0.387
## 8 ARIMA210   Test      0.997 1.47 1.03 27.5 29.8 0.923 0.954 0.730
## 9 ARIMA011   Test      1.05 1.51 1.06 30.3 30.8 0.953 0.982 0.731
## 10 arimaBIC   Test      1.05 1.51 1.06 30.3 30.8 0.953 0.982 0.731
## 11 arimaAICc  Test      1.39 1.90 1.39 43.3 43.5 1.25 1.23 0.734
## 12 etsAICc    Test      1.53 2.26 1.53 39.9 39.9 1.37 1.46 0.807
## 13 naive      Test      1.68 2.39 1.68 47.2 47.2 1.51 1.55 0.806
## 14 etsBIC     Test      1.68 2.39 1.68 47.2 47.2 1.51 1.55 0.806
```

```
HCV.fit %>%
  forecast(h = 12) %>%
  autoplot(CPI)
```

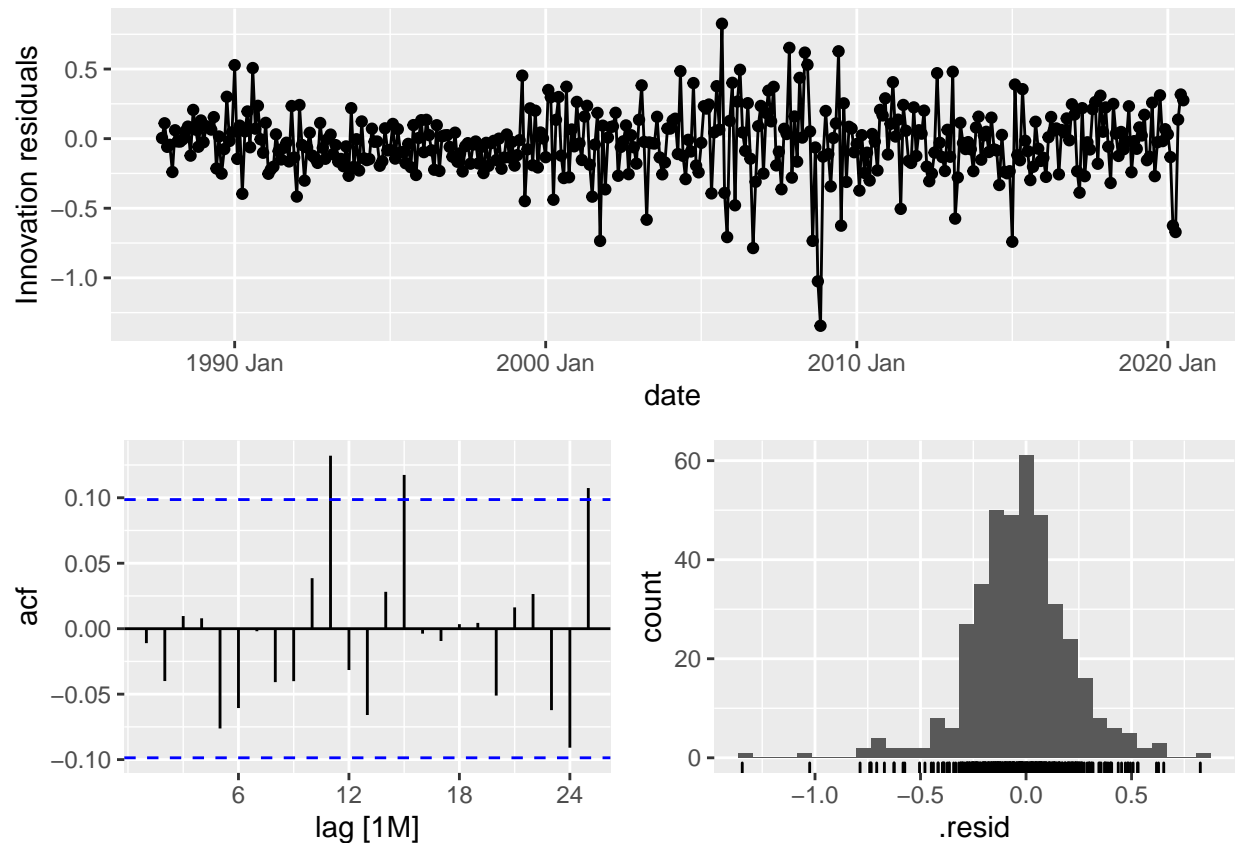


The best model according to RMSE in the test set is the ARIMA(2,1,0)(0,0,2)[12] while in the training set it is the ARIMA(2,1,0)(2,0,1)[12] w/drift. The prior has a better mix of minimum AICc and BIC so it will be selected as the best model

```
HCV.fit %>% augment %>% features(.innov, ljung_box)
```

```
## # A tibble: 7 x 3
##   .model      lb_stat lb_pvalue
##   <chr>      <dbl>    <dbl>
## 1 ARIMA011    0.391    5.32e- 1
## 2 ARIMA210    0.0485   8.26e- 1
## 3 arimaAICc   0.0000288 9.96e- 1
## 4 arimaBIC    0.391    5.32e- 1
## 5 etsAICc    24.7      6.80e- 7
## 6 etsBIC     59.6     1.18e-14
## 7 naive      59.4     1.29e-14
```


```
HCV.fit %>% select(ARIMA210) %>% gg_tsresiduals()
```



All ARIMA models appear to have residuals similar to white noise according to the Ljung-Box test.

Forecasting

```
CPI %>%
  model(arima = ARIMA(pc1_cpi ~ pdq(2,1,0) + PDQ(0,0,2)),
        ets = ETS(pc1_cpi, ic = "aicc"),
        naive = NAIVE(pc1_cpi)) %>%
  forecast(h = 12) %>%
  autoplot(CPI)
```



```
CPI %>%
  model(arima = ARIMA(pc1_cpi ~ pdq(2,1,0) + PDQ(0,0,2)),
        ets = ETS(pc1_cpi, ic = "aicc"),
        naive = NAIVE(pc1_cpi)) %>%
  forecast(h = 1)
```

```
## # A tibble: 3 x 4 [1M]
## # Key:   .model [3]
##   .model    date      pc1_cpi .mean
##   <chr>    <mth>      <dist> <dbl>
## 1 arima  2021 Aug N(5.3, 0.061)  5.30
## 2 ets    2021 Aug N(5.6, 0.14)  5.56
## 3 naive  2021 Aug N(5.4, 0.15)  5.37
```

The ARIMA model produces a point forecast of 5.30 for August 2021.

Official Results

The Bureau of Labor Statistics released the official inflation measures for the month of August 2021 on September 14, 2021 at 8:30am. Rounded to 2 decimal places, official year over year inflation in August was 5.25 compared to my forecasted figure of 5.30.